

Inf2C Computer Systems

Coursework 2

MIPS Processor Simulator

Deadline: Wed, 27 Nov, 16:00

Instructor: Boris Grot

TA: Siavash Katebzadeh

The aim of this assignment is to write a simulator for a 5-stage multi-cycle MIPS processor with a simple, direct-mapped cache. A simulator is nothing more than a functional model of a processor that mimics the behavior of a real processor but is written in a high level language. Your simulator, written in C, will read a memory file consisting of MIPS instructions and data, “execute” the instructions and output the state of the processor and various statistics for the cache during the execution. To get you started, you will be provided with a skeleton implementation of the simulator that you will need to extend. You are strongly advised to read up on MIPS processor and caches in the lecture notes and course textbook and to commence work as soon as possible.

This is the second of the two assignments for the Inf2C-CS course. It is worth 50% of the coursework marks and 20% of the overall course marks.

Please bear in mind that the guidelines on academic misconduct from the Undergraduate year 2 student handbook are available on the following link <http://web.inf.ed.ac.uk/infweb/student-services/ito/students/year2>.

1 Overview

In this assignment, you are provided with a skeleton of the simulator, which includes the 5-stage MIPS processor and memory. You will need to extend the provided code to support a set of specified MIPS instructions (Task 1) and add a cache between the processor and the memory (Task 2).

Skeleton Organization: The skeleton is broken down into four source code files. Each file accomplishes particular tasks. You are allowed to modify and submit only certain files. The functionality and permission to modify for each file are described in Table 1.

Filename	Functionality	Modifiable
<code>mipssim.c</code>	Multi-cycle MIPS processor (datapath + control)	Yes
<code>mipssim.h</code>	Data structure definitions for datapath	No
<code>memory_hierarchy.c</code>	Memory hierarchy implementation	Yes
<code>parser.h</code>	Reading and parsing input files	No

Table 1: Source Files

mipssim.c: This file describes the multicycle MIPS processor as studied in class. The processor consists of the following core components: PC, Pipeline registers (IR, A, B, MDR, and ALUOut), Programmer-visible registers (in the register file), ALU, ALU Control, and Control.

The processor’s functionality can be broken down into the following logical stages: *instruction_fetch*, *decode_and_read_RF*, *execute*, *memory_access* and *write_back*. Note that each stage updates some architectural or microarchitectural state. For instance, the *instruction_fetch* stage updates the IR. The *write_back* stage updates the RF.

Furthermore, this file handles the Control component by implementing a finite state machine. The state machine can be found in a function named *FSM*.

mipssim.h: This file defines the following required data structures for the MIPS processor:

- *ctrl_signals*, which control the datapath and are updated on a cycle-by-cycle basis by the control FSM
- *instr_meta*, which stores information about the instruction currently stored in the IR
- *memory_stats_t*, which consists of memory stats for loads, stores and instruction fetches
- *pipe_regs*, which includes the PC and microarchitectural registers of the processor (IR, A, B, ALUOut and MDR). For convenience, we refer to these registers as *pipeline registers* to indicate that they are spread out over the datapath (unlike the programmer-visible registers, which are all located inside the register file).

On any given cycle, the complete state of the processor is stored in a structure called *architectural_state*. This structure includes the current clock cycle since the start of execution (*clock_cycle*), state of the current instruction’s execution (e.g., INSTR_FETCH or DECODE), current values of control signals (*control*) and memory stats (*mem_stats*). This structure also includes an array of *registers*, which models a register file, as well as the *memory*. Finally, *architectural_state* also includes the pipeline registers, which are maintained in two *pipe_regs* structs: *curr_pipe_regs* and *next_pipe_regs*. The former

2 Task 1: MIPS Processor

In the first task, your job is to complete the datapath of a MIPS processor provided in the skeleton *mipssim.c*. Figure 1 illustrates the processor datapath and the control signals. In the skeleton, there are five functions: *instruction_fetch*, *decode_and_read_RF*, *execute*, *memory_access* and *write_back*, corresponding to the five stages of the processor. Some of these functions are incomplete, and you need to implement them according to the MIPS data and control paths as specified in the P&H chapter on the multi-cycle processor (available on Learn).

Additionally, you are required to complete the Control implementation in *mipssim.c*. In every cycle, the Control component controls the state of each stage of the processor using a finite state machine (FSM). Figure 2 illustrates the FSM used in the Control component of the multi-cycle MIPS processor as discussed in class. Note that the names of the states in Figure 2 may be different from the source files provided to you. Each circle in Figure 2 corresponds to one state of the FSM and lists (1) the values for mux select signals that are set in that state, and (2) all enable signals that are set in that state. For instance, State 0 (i.e. *instruction fetch*) shows that enable signals *MemRead*, *IRWrite* and *PCWrite* must be set.

In the skeleton, a global struct named *arch_state.control* contains the control signals generated by the Control component. You **must use these signals** and complete the *FSM* function to control the 5 stages. You may add new FSM states if necessary.

The skeleton already implements both the datapath and the control for the *ADD* instruction. Your task is to extend the code to support the following instructions: *LW*, *SW*, *ADDI*, *J*, *BQE* and *SLT*. Note that Figure 2 does not show the required states and transitions for an *ADDI* instruction. You must design the state machine for *ADDI* by yourself (**Hint:** *ADDI* is similar to *ADD* except it has one different operand).

3 Task 2: Cache

In this task, your job is to extend your simulator and add a cache to the memory hierarchy. The structure of the cache must be direct-mapped, which means that each memory address maps to exactly one location in the cache. For each cache block, you need to store the data, a tag and a valid bit.

You must implement the cache in the *memory_hierarchy.c* file. In this file, there are two memory functions already defined: *memory_read* and *memory_write*. The *memory_read* function is used to fetch instructions and load data from memory. The *memory_write* function writes data to the memory. You must extend these functions to access the cache. You must also update the relevant cache hits statistic, as well as the content of the cache, as necessary. The statistics are maintained in a global structure *arch_state.mem_stats*, which is updated on every access to the memory hierarchy.

The cache will use the following parameters:

Fixed parameters:

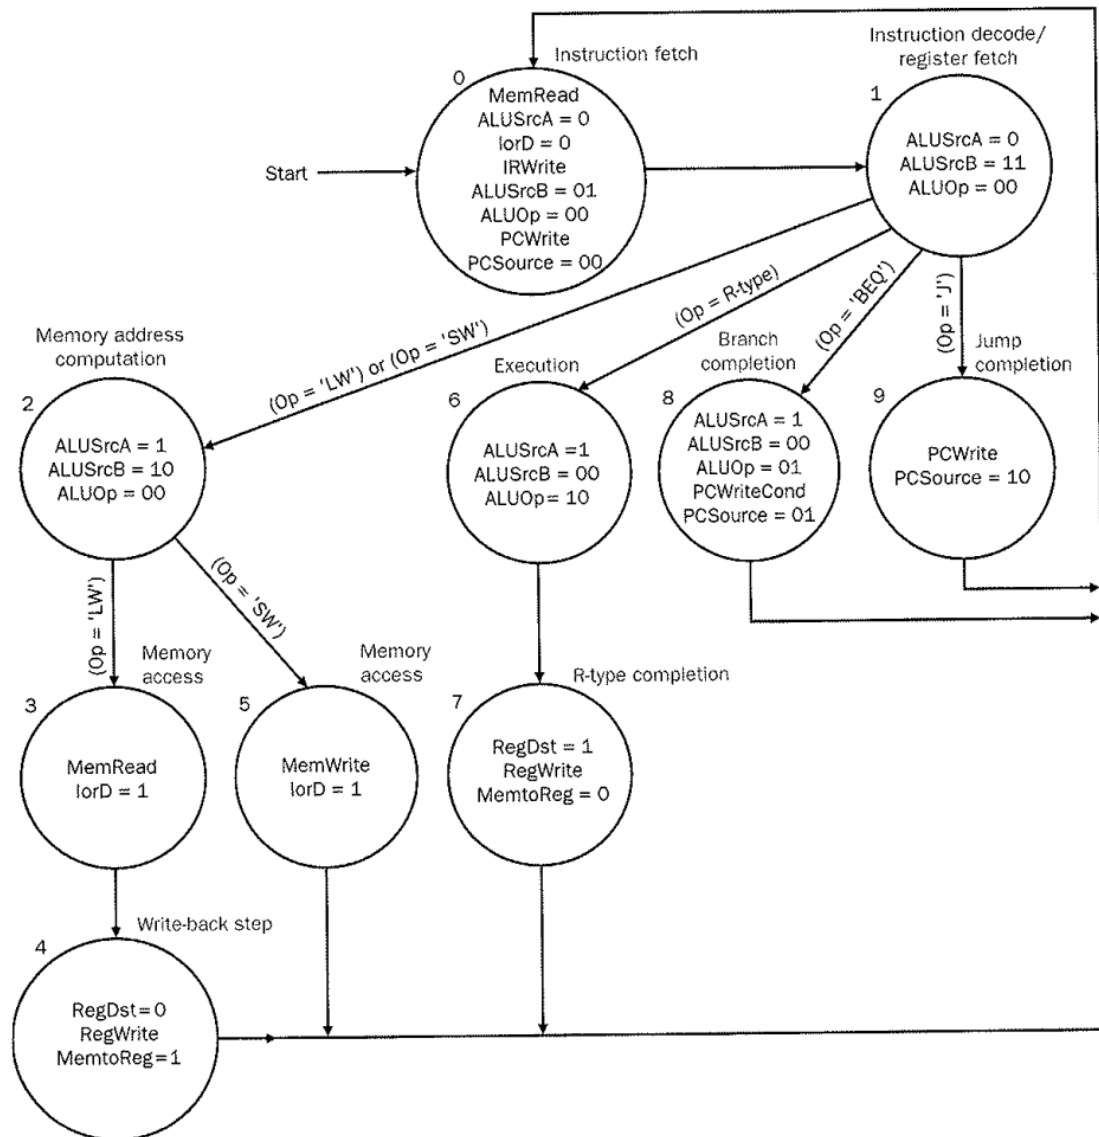


Figure 2: Finite state machine used for a multi-cycle MIPS processor

1. Addresses are 32-bits
2. Memory is byte addressable
3. Cache block size is 16 bytes
4. Cache holds both instructions and data (this is called a *unified* cache)
5. Cache uses a write-through policy
6. Cache uses a write-no-allocate policy; i.e., if a store does not find the target block in the cache, it does not allocate it.

Variable parameter: The size of the cache (in bytes) is passed as a command-line argument to the program and stored for you in a global variable named *cache_size*. Note that this parameter specifies the size of the *data* portion of the cache. Any other information that the cache needs to store (e.g., tags) are additional to this.

Notes:

- Instruction fetch uses the *memory_read* function. Hence, each instruction fetch updates the cache hit statistics just like loads.
- If the cache size is set to **0**, the cache is disabled and the processor directly accesses main memory.
- The cache will not be larger than 16KB.
- You **must** use dynamic memory allocation (*malloc*) for your cache structure.

Based on the cache size, you need to calculate the number of bits required for the tag field, which will be stored in *arch_state.bits_for_cache_tag*. You must set this variable by providing a correct value in the call to the *memory_stats_init* function in *memory_hierarchy.c*. The function itself is already defined for you.

Summary: In this assignment, you are required to complete a 5-stage multi-cycle MIPS simulator which includes a cache model. You will need to fill in the control FSM for instructions *LW*, *SW*, *ADDI*, *J*, *BQE* and *SLT*. Your FSM must use the provided control signals (in *arch_state.control*). You will also need to correctly maintain the pipeline registers (*arch_state.curr_pipe_regs*), programmer-visible registers (*arch_state.registers*), and memory (*arch_state.memory*).

Furthermore, you will need to appropriately create and configure the cache data structure(s) based on parameters that are passed as command-line arguments. You will need to set the *arch_state.bits_for_cache_tag* variable to the number of bits for the tag via a call to *memory_stats_init* and update the cache hit statistics (in *arch_state.mem_stats*) as necessary for both instruction fetches and data accesses.

4 Notes on the Implementation

1. You are expected to dynamically allocate memory using **malloc** for the cache data structure. Submissions that use statically allocated memory (e.g., statically declared arrays of fixed size) will be penalized. Variable-length arrays that can be sized based on a runtime param in C99 are **NOT** allowed.
2. There are many ways to implement the cache structure. The details of the implementation are entirely up to you – you are the designer! The only thing that matters is correctness at the functional level. Just remember, you are coding in a high level language – think data structures, not bits and gates.

3. You can use the C library functions available from the header files that are already included in `mipssim.c` and `mipssim.h`. You may **not** use any library functions beyond these (i.e., do not include other C header files).
4. You will not be marked on how fast your simulator runs or how much memory it uses. The only criteria for evaluating your implementation of both the MIPS processor and the cache is correctness.

5 Input Files

The skeleton will read two files:

1. *The memory state file*: a text file, which is a mix of instructions and data represented by a sequence of 0 and 1 characters, one word (32 bits) per line. Note that the *first* non-comment line of the memory file is always an instruction, which starts at address 0x0. Subsequent words are placed in consecutive memory locations. For this assignment, a special instruction with opcode `111111` (in binary) is considered as the End-Of-Program (EOP) instruction, which terminates the program. In the provided skeleton, there are two memory state file examples: `memfile-simple.txt` and `memfile-complex.txt`.
2. *The register state file*: a text file, which contains the initial state of programmer-visible registers. This file has up to 31 uncommented lines for registers (i.e., all of the programmer-visible registers except \$0) starting with register \$1. Each line specifies a decimal value to which the corresponding register will be initialized. If fewer than 31 values are specified, the remaining registers will be initialized with zeros. In the provided skeleton, there is one register state file example: `regfile.txt`.

6 Output Format

You do not need to generate any output for marking purposes. Instead, you must ensure that in each cycle all relevant variables/structs of the `arch_state` struct are properly updated. The automated marking will check the `arch_state` struct in each cycle, checking the correctness of the control signals (i.e., the `arch_state.control` field), datapath state (i.e., fields: `arch_state.curr_pipe_regs` and `arch_state.registers`) and memory state (i.e. fields: `arch_state.mem_stats`, `arch_state.bits_for_cache_tag`, `arch_state.memory`). The automated marking will use the functions `marking_after_clock_cycle()` and `marking_at_the_end()`. Make sure that you **do not use or modify** these functions!

7 Debugging

1. To verify correctness of your implementation, you should write your own memory and register state files. For verifying the datapath, start testing each instruction

individually. For verifying the cache, you need to come up with test cases to predictably generate certain behaviors (hits and misses). For instance, think of a trace with four accesses that uses four different addresses and has a 50% cache hit rate in a direct-mapped cache.

2. Your simulator will be tested with memory and register state files different from the provided ones.
3. You can use *printf* function for debugging your code. It will not affect your marking.

8 Compiling and Running the Simulator

You **must** compile the simulator on the DICE machines with the following command:

```
gcc -o mipssim mipssim.c memory_hierarchy.c -std=gnu99 -lm
```

Note that this is the exact command we will use for compiling your code for marking purposes. Compiling the source files creates an executable **mipssim**. Make sure that your simulator both compiles with the exact command and runs on a **DICE** machine without errors and warnings. Otherwise, you will receive a **0** mark.

The following are examples of invoking the simulator with valid command-line parameters.

```
./mipssim 128 memfile-simple.txt regfile.txt
```

Where **mipssim** is the name of the executable file, **128** indicates cache is enabled and total size of cache is 128 bytes, **memfile-simple.txt** is the name of the memory file and **regfile.txt** is the name of the register state file. Both memory and register state files are located in the same directory as **mipssim**. Another example is:

```
./mipssim 0 memfile-simple.txt regfile.txt
```

Where **0** indicates that the cache is disabled.

9 Submission

1. You should submit a copy of your simulator by 4pm on November 27, 2019 using the following command at a command-line prompt on a DICE machine.

```
submit inf2c-cs cw2 mipssim.c memory_hierarchy.c
```


2. You can submit only `mipssim.c` and `memory_hierarchy.c` files. Please put all of your definitions and implementations inside of these two files.
3. Unless there are special circumstances, **late submissions are not allowed and will receive a mark of 0**. Please consult the online undergraduate year 2 student handbook for further information on this.
4. You can submit more than once up until the submission deadline. All submissions are timestamped automatically. Identically named files will overwrite earlier submitted versions, so we will mark the latest submission that comes in before the deadline.

Warning: Unfortunately the `submit` command will technically allow you to submit late even if you submitted before the deadline. **Don't do this!** We can only retrieve the latest version, which means you will receive a **0** mark for submitting after the submission deadline.

For additional information about late penalties and extension requests, see the School web page below. Do **NOT** email any course staff directly about extension requests; you must follow the instructions on the web page.

<http://web.inf.ed.ac.uk/infweb/student-services/ito/admin/coursework-projects/late-coursework-extension-requests>.

10 Assessment

The assignment will be auto-marked. Task 1 and task 2 are worth 50% each. Your solutions will be evaluated with a number of test inputs. For each task, your mark will be proportional to the pass rate of the tests.

11 Similarity Checking and Academic Misconduct

You must submit your own work. Any code that is not written by you must be clearly identified and explained through comments at the top of your files. Failure to do so is plagiarism. Detailed guidelines on what constitutes plagiarism can be found at: <http://web.inf.ed.ac.uk/infweb/admin/policies/guidelines-plagiarism>. All submitted code is checked for similarity with other submissions using the MOSS¹ system. MOSS has been effective in the past at finding similarities. It is not fooled by name changes and reordering of code blocks.

12 Questions

If you have any questions about the assignment, please start by checking existing discussions on Piazza – chances are, others have already encountered (and, possibly,

¹<http://theory.stanford.edu/~aiken/moss/>

solved) the same problem. If you can't find the answer to your question, start a new discussion. You should also take advantage of the drop-in labs and the lab demonstrators who are there to answer your questions.

November 12, 2019